



[◀ Back to Deep Learning](#)

Dog Breed Classifier

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Congratulations for passing this project!
Good luck with future submissions :)

Files Submitted

The submission includes all required files.

Awesome! The submission includes all the required files.

Step 1: Detect Humans

The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected human face.

Perfect!

Percentage of human faces detected in human images = 100 %

Percentage of human faces detected in dog images = 11 %

The submission opines whether Haar cascades for face detection are an appropriate technique for human detection.

Good explanation.

Some issues pertaining to face detection using Haar Cascades approach can be seen here: [1](#), [2](#).

Step 2: Detect Dogs

The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected dog.

The percentage values are even better here :)

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

The submission specifies a CNN architecture.

Good job building a CNN from scratch.

Here you can also add `batchNormalization()` layer to improve model performance. See the documentation [here](#).

Also check [this](#) post by Ian Goodfellow.

The submission specifies the number of epochs used to train the algorithm.

You can use an [early stopping callback](#) provided by Keras, which ends training based on pre configured rules such as non improving accuracy.

The trained model attains at least 1% accuracy on the test set.

The accuracy is 1.5550 %.

Step 5: Create a CNN to Classify Dog Breeds

The submission downloads the bottleneck features corresponding to one of the Keras pre-trained models (VGG-19, ResNet-50, Inception, or Xception).

Awesome!
The bottleneck features are successfully imported.

The submission specifies a model architecture.

The submission details why the chosen architecture succeeded in the classification task and why earlier attempts were not as successful.

The submission compiles the architecture by specifying the loss function and optimizer.

Nice job using `categorical_crossentropy` and `rmsprop`. For a detailed overview of gradient descent optimizers check out this [link](#).

The submission uses model checkpointing to train the model and saves the model weights with the best validation loss.

Perfect!

The submission loads the model weights that attained the least validation loss.

Your use of `save_best_only=True` suggests the model with least validation loss is loaded.

Accuracy on the test set is 60% or greater.

The accuracy is 78.3493 %.

The submission includes a function that takes a file path to an image as input and returns the dog breed that is predicted by the CNN.

Well done :)

Step 6: Write Your Algorithm

The submission uses the CNN from Step 5 to detect dog breed. The submission has different output for each detected image type (dog, human, other) and provides either predicted actual (or resembling) dog breed.

Great! Different outputs are generated for each detected image :)

Step 7: Test Your Algorithm

The submission tests at least 6 images, including at least two human and two dog images.

Nice improvement suggestions! Check this [blog](#).

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)

[Rate this review](#)

